

Managing data reliability and integrity in federated cloud storage

Krzysztof Styrz

AGH University of Science and Technology, Kraków
kstyrz@gmail.com



Supervisor: Marian Bubak, PhD
Advisor: Piotr Nowakowski

Agenda

- 1 Background
 - Objectives
 - Motivation
- 2 Design
 - VPH-Share
 - Architecture
- 3 Validation
 - Cloud storage
 - Algorithm
 - Technologies
 - Efficiency
- 4 Approaches
 - Classic
 - PoR
 - DIP
 - Solution
- 5 Conclusions

Objectives

Data in VPH-Share Cloud Platform:

- **sensitive** data on one or more cloud storage nodes (S3, Swift, ...),
- data has to be **available** and **valid**.

The aims of this work:

- ① **efficient** data validation mechanism in federated cloud storage,
- ② implementation under various practical **limitations**.

Essentials aspects of this work:

- design and implementation of DRI REST WS enabling data validation,
- integration with the VPH-Share Cloud Platform environment,
- design of a bandwidth-efficient validation algorithm

Data integrity in cloud storage

Why data validation in cloud storage?

SLA guarantees 99.9%, if not met – credit return:

- <http://aws.amazon.com/s3-sla/>
- <http://www.rackspace.com/cloud/legal/sla/>
- <https://developers.google.com/appengine/sla>







Moreover, in replicated environment one can use other cloud provider while primary is unavailable or contains corrupted data.

Recent problems with the *cloud* services:


- Gmail: erased emails, blocked accounts,
- Amazon S3: serving corrupted data, unavailable data,
- GoogleDocs: unauthorized data access,

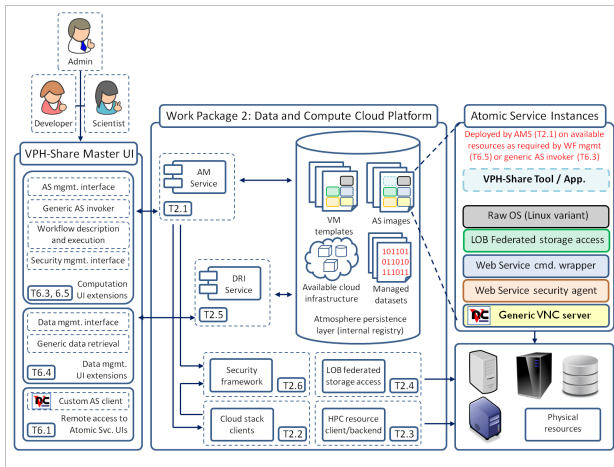
VPH-Share environment

- **Atmosphere Internal Registry (AIR)** – metadata database of datasets, files, storage nodes, configuration, policy,
- data structure: *dataset = set of files*



Storage nodes – direct access to cloud storage nodes rather than LOBCDER layer (S3, Swift, ...).





Functional requirements:

- periodic and on-request validation,
- data replication,
- user notification about integrity errors.


Nonfunctional requirements:

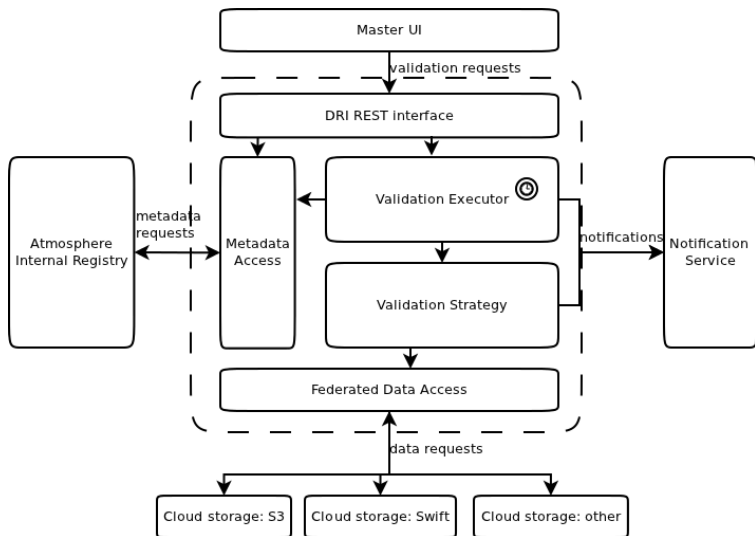
- **bandwidth-efficient** validation mechanism,
- scalability,
- configurability.

DRI architecture

DRIService
+ registerDataset(dataset : ManagedDatasetDescription) : ManagedDatasetID
+ unregisterDataset(id : ManagedDatasetID)
+ replicateDatasetToResource(id : ManagedDatasetID, source : DataSourceID)
+ dereplicateDatasetFromResource(id : ManagedDatasetID, source : DataSourceID)
+ datasetChanged(id : ManagedDatasetID, dataset : ManagedDatasetDescription)
+ validateDataset(id : ManagedDatasetID) : Message
+ setManagementPolicy(policy : ManagementPolicy)
+ getManagementPolicy(id : ManagedDatasetID) : ManagementPolicy

DRI Service





Cloud storage characteristics

- hierarchical structure: *account / container / object*,
- typical operations: *list, get, put, delete, update, ...*,
- REST API goes toward CDMI/OCCI standards:
 - <http://docs.amazonwebservices.com/AmazonS3/latest/API/>
 - <http://docs.openstack.org/api/openstack-object-storage/1.0/content/>
 - <http://docs.rackspace.com/files/api/v1/cf-devguide/content/index.html>
- SLA quality, best quality–price ratio, scalability, ... ,

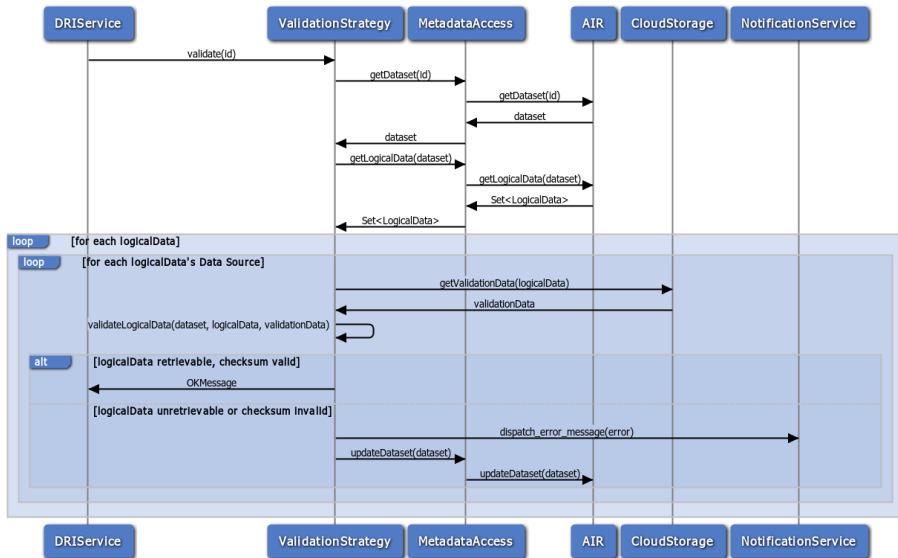
No support for HTTP Multi-Range requests:

- new request for every chunk,
- flooding the cloud storage,
- really bad performance → RTT for every chunk.

Validation algorithm pseudocode

```
1 def validate_dataset(dataset id) {  
2     dataset = get_dataset_metadata(id);  
3     files = get_dataset_files(dataset);  
4     for file in files {  
5         for data_source in dataset.get_data_sources() {  
6             data = get_validation_data(file, data_source);  
7             if data == null # unable to retrieve data  
8                 dispatch_error_message(file, data_source);  
9             result = validate_data(data, file, dataset);  
10            if not result # file is invalid  
11                dispatch_error_message(file, data_source);  
12        }  
13    }  
14 }
```

Validation algorithm sequence diagram



Implementation technologies

- **Jersey API** – JAX-RS reference implementation, REST communication
- **JClouds** – very good cloud operability library,
- **Quartz** – tasks execution and scheduling library,
- **Guice** – dependency injection and modularity library,
- **Maven** – build + dependencies tool,
- **Apache Tomcat** – server.

Validation efficiency

Significant size of data stored in federated cloud storage makes whole-file checksum-based validation completely inefficient. Additionally, occupying substantial portion of network bandwidth is highly undesirable.

For example:

typical dataset ≈ 10 GB \rightarrow goal: reduce to $\approx 10\%$

Can we do better?

Generally yes, but lots of practical limitations:

- cloud storage limitations,
- files cannot be modified in VPH-Share.

Promising approaches:


- 1 Proofs of Retrievability (PoR),
- 2 Data integrity proofs (DIP).

Classical integrity mechanisms

- **Hash functions** – for detecting data corruption (MD5, SHA1, ...),
- **Error Correcting Codes (ECC)** – for correcting small parts of data,
- **Message Authentication Codes (MAC)** – hash + auth_key for secure authentication.

Validation mechanism steps:

- 1 setup: compute and store,
- 2 validate: compute + compare.



Proofs of Retrievability scheme

Setup:

- 1 divide into chunks,
- 2 compute ECC + MAC,
- 3 $\frac{F^*}{F} \approx 102 - 115\%$.

Validation:


- 1 challenge-response rounds:
- 2 more rounds \rightarrow higher probability,
- 3 more rounds \rightarrow higher bandwidth.

Problems:


- 1 we cannot modify stored file,
- 2 data is served by LOBCDER, ECC useless.

ECC + MAC: Verification

Archive



F^*



Data integrity proofs scheme

Setup:


- ➊ divide into k chunks,
- ➋ pick m bits in every chunk,
- ➌ compute MAC and store.

Validation:

- ➊ compute as in setup,
- ➋ compare with original MAC.

Problems:

- ➊ cannot query cloud storage bit by bit,
- ➋ cannot append to stored file.



Our solution – mixed PoR + DIP

Setup:



- 1 divide into chunks of size S ,
- 2 compute MAC of every single chunk and store.

Validation:

- 1 pick sequence of k random indexes (Mersenne-Twister),
- 2 get k chunks data,
- 3 compute and verify with originals.

Solution features:

- configurable accuracy–overhead ratio,
- practical feasibility.



Work status:

- working periodic and on-request validation,
- integrity errors notifications page,
- task scheduling + configurability,
- **testing the implementation of our solution,**
- **thesis writing start.**

Conclusions:

- cloud storages have their usability limitations,
- remote data validation – not a trivial task,
- efficient validation only with high probability.

Future work:

- integrate LOBCDER and DRI tools,
- watch future versions of CDMI/OCCI standards for significant improvements.